
Kalman Filter Theory and Applications

Heuristic Overview

Michael L. Carroll

July 27, 2020

©2020 by Michael L. Carroll

Part I

The Basic Kalman Equations Topics

Part I

The Basic Kalman Equations

Topics

- Understanding the Equations: Heuristic Introduction

Part I

The Basic Kalman Equations

Topics

- Understanding the Equations: Heuristic Introduction
- Equation Drilldown: Taking the Equations Apart

Part I

The Basic Kalman Equations

Topics

- Understanding the Equations: Heuristic Introduction
- Equation Drilldown: Taking the Equations Apart
- State Space Concepts

Part I

The Basic Kalman Equations

Topics

- Understanding the Equations: Heuristic Introduction
 - Equation Drilldown: Taking the Equations Apart
 - State Space Concepts
-

The Basic Kalman Equations

Understanding the Equations: Heuristic Introduction

Subtopics

The Basic Kalman Equations

Understanding the Equations: Heuristic Introduction

Subtopics

- Recursive Predictor-Corrector Algorithms

The Basic Kalman Equations

Understanding the Equations: Heuristic Introduction

Subtopics

- Recursive Predictor-Corrector Algorithms
- **Running Averages (Part B)**

The Basic Kalman Equations

Understanding the Equations: Heuristic Introduction

Subtopics

- Recursive Predictor-Corrector Algorithms
- **Running Averages (Part B)**
- Prediction (Extrapolation)

The Basic Kalman Equations

Understanding the Equations: Heuristic Introduction

Subtopics

- Recursive Predictor-Corrector Algorithms
- **Running Averages (Part B)**
- Prediction (Extrapolation)
- Correction (Measurement Update)

The Basic Kalman Equations

Understanding the Equations: Heuristic Introduction

Subtopics

- Recursive Predictor-Corrector Algorithms
- **Running Averages (Part B)**
- Prediction (Extrapolation)
- Correction (Measurement Update)
- Gain Computation

The Basic Kalman Equations

Understanding the Equations: Heuristic Introduction


Subtopics

- Recursive Predictor-Corrector Algorithms
 - **Running Averages (Part B)**
 - Prediction (Extrapolation)
 - Correction (Measurement Update)
 - Gain Computation
-

Recursive Running Average Python Scripts:

`dyn_simple_moving_avg_demo.py`

Recursive Running Average Python Scripts: dyn_simple_moving_avg_demo.py



Running Average

Shortcomings

Running Average

Shortcomings

- Uses (implicitly) all data points, even those from the "remote" past

Running Average

Shortcomings

- Uses (implicitly) all data points, even those from the "remote" past
- Gain gets smaller and smaller ... running average filter goes to sleep, becomes insensitive to new information

Running Average

Shortcomings

- Uses (implicitly) all data points, even those from the "remote" past
- Gain gets smaller and smaller ... running average filter goes to sleep, becomes insensitive to new information
- Can improve things with a simple moving average (sliding window)

Running Average

Shortcomings

- Uses (implicitly) all data points, even those from the "remote" past
 - Gain gets smaller and smaller ... running average filter goes to sleep, becomes insensitive to new information
 - Can improve things with a simple moving average (sliding window)
-

Simple Moving Average

Addressing the Shotcomings

Simple Moving Average

Addressing the Shotcomings

- Simple Moving Average of Window Size M only averages the M most recent measurements

Simple Moving Average

Addressing the Shotcomings

- Simple Moving Average of Window Size M only averages the M most recent measurements
- Has a fixed gain of $1/M$ which never goes to zero

Simple Moving Average

Addressing the Shotcomings

- Simple Moving Average of Window Size M only averages the M most recent measurements
- Has a fixed gain of $1/M$ which never goes to zero
- Assume the number of available data points is $N > M$

Simple Moving Average

Addressing the Shotcomings

- Simple Moving Average of Window Size M only averages the M most recent measurements
- Has a fixed gain of $1/M$ which never goes to zero
- Assume the number of available data points is $N > M$
- Simple Moving Average of window size M is
$$s(k, M) = \frac{1}{M} \sum_{j=k-M+1}^k z(j), \text{ valid for } k \geq M - 1 \text{ if indexing starts at } 0$$
(we can set the SMA equal to the running average for $k < M - 1$)

Simple Moving Average

Addressing the Shotcomings

- Simple Moving Average of Window Size M only averages the M most recent measurements
- Has a fixed gain of $1/M$ which never goes to zero
- Assume the number of available data points is $N > M$
- Simple Moving Average of window size M is
$$s(k, M) = \frac{1}{M} \sum_{j=k-M+1}^k z(j), \text{ valid for } k \geq M - 1 \text{ if indexing starts at } 0 \text{ (we can set the SMA equal to the running average for } k < M - 1)$$
- Recursive SMA formula:
$$s(k, M) := s(k - 1, M) + \frac{1}{M} [z(k) - z(k - M + 1)]$$

Simple Moving Average

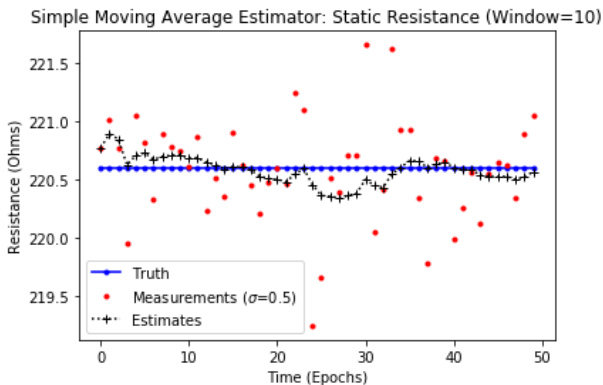
Addressing the Shotcomings

- Simple Moving Average of Window Size M only averages the M most recent measurements
- Has a fixed gain of $1/M$ which never goes to zero
- Assume the number of available data points is $N > M$
- Simple Moving Average of window size M is
$$s(k, M) = \frac{1}{M} \sum_{j=k-M+1}^k z(j), \text{ valid for } k \geq M - 1 \text{ if indexing starts at } 0 \text{ (we can set the SMA equal to the running average for } k < M - 1)$$
- Recursive SMA formula:
$$s(k, M) := s(k - 1, M) + \frac{1}{M} [z(k) - z(k - M + 1)]$$

Simple Moving Average Python Scripts:

`simple_moving_avg_demo.py` of window size 10

Simple Moving Average Python Scripts: simple_moving_avg_demo.py of window size 10



Running Average

Addressing the Shotcomings

Running Average

Addressing the Shortcomings

- Note that the recursive SMA seems to have shortcomings of its own

Running Average

Addressing the Shortcomings

- Note that the recursive SMA seems to have shortcomings of its own
- First, it seems overly conservative about updating the previous SMA value

Running Average

Addressing the Shotcomings

- Note that the recursive SMA seems to have shortcomings of its own
- First, it seems overly conservative about updating the previous SMA value
 - However, this is because unchanging process dynamics means that measurements do not change much

Running Average

Addressing the Shortcomings

- Note that the recursive SMA seems to have shortcomings of its own
- First, it seems overly conservative about updating the previous SMA value
 - However, this is because unchanging process dynamics means that measurements do not change much
 - Performs worse than running average on constant dynamics, but better with changing dynamics (See next slide)

Running Average

Addressing the Shortcomings

- Note that the recursive SMA seems to have shortcomings of its own
- First, it seems overly conservative about updating the previous SMA value
 - However, this is because unchanging process dynamics means that measurements do not change much
 - Performs worse than running average on constant dynamics, but better with changing dynamics (See next slide)
- A more aggressive formulation of the recursive SMA can be defined by using a weighting scheme whereby the oldest measurement is weighted less than the most recent

Running Average

Addressing the Shortcomings

- Note that the recursive SMA seems to have shortcomings of its own
- First, it seems overly conservative about updating the previous SMA value
 - However, this is because unchanging process dynamics means that measurements do not change much
 - Performs worse than running average on constant dynamics, but better with changing dynamics (See next slide)
- A more aggressive formulation of the recursive SMA can be defined by using a weighting scheme whereby the oldest measurement is weighted less than the most recent
- Neither of these averaging filters explicitly takes into account the degree of uncertainty in the process dynamics

Running Average

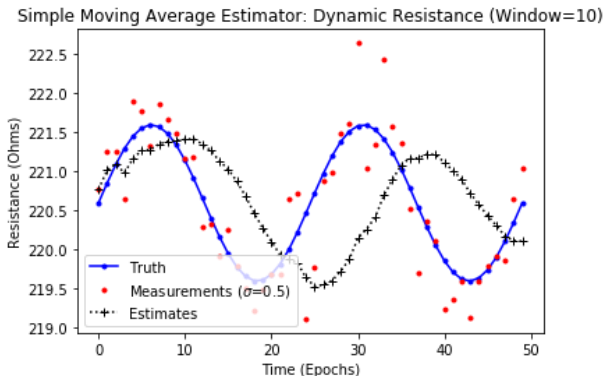
Addressing the Shortcomings

- Note that the recursive SMA seems to have shortcomings of its own
 - First, it seems overly conservative about updating the previous SMA value
 - However, this is because unchanging process dynamics means that measurements do not change much
 - Performs worse than running average on constant dynamics, but better with changing dynamics (See next slide)
 - A more aggressive formulation of the recursive SMA can be defined by using a weighting scheme whereby the oldest measurement is weighted less than the most recent
 - Neither of these averaging filters explicitly takes into account the degree of uncertainty in the process dynamics
-

Simple Moving Average Python Scripts:
`dyn_simple_moving_avg_demo.py` of window size 10

Simple Moving Average Python Scripts:

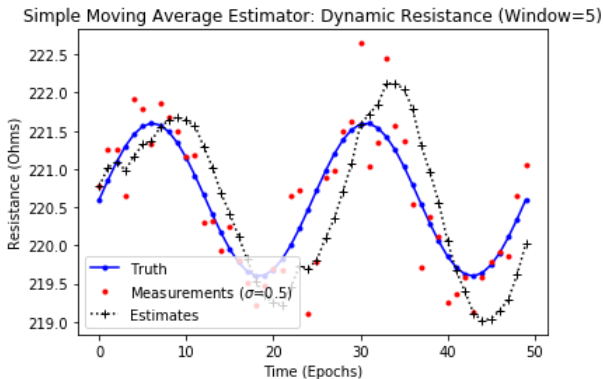
`dyn_simple_moving_avg_demo.py` of window size 10



Simple Moving Average Python Scripts:
`dyn_simple_moving_avg_demo.py` of window size 5

Simple Moving Average Python Scripts:

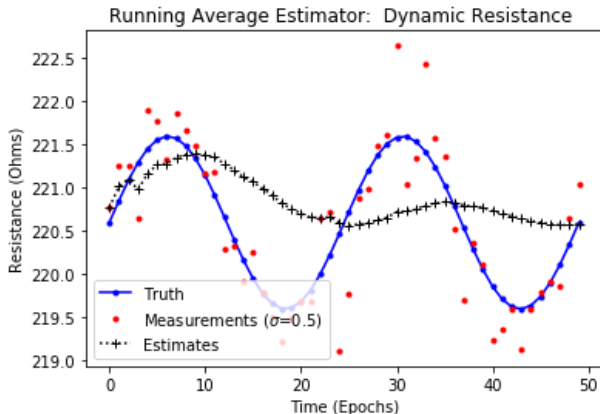
`dyn_simple_moving_avg_demo.py` of window size 5



Simple Moving Average Python Scripts:
`dyn_simple_moving_avg_demo.py` of window size 5

Simple Moving Average Python Scripts:

`dyn_simple_moving_avg_demo.py` of window size 5



Running Average Summary

Vector Check

Running Average Summary

Vector Check

- Where are we?

Running Average Summary

Vector Check

- Where are we?
 - Looked at some Python implementations of the running average and simple moving average

Running Average Summary

Vector Check

- Where are we?
 - Looked at some Python implementations of the running average and simple moving average
 - Ran some simulations of both constant and dynamic scalar processes and examine how the two types of averages performed

Running Average Summary

Vector Check

- Where are we?
 - Looked at some Python implementations of the running average and simple moving average
 - Ran some simulations of both constant and dynamic scalar processes and examine how the two types of averages performed
- What's next?

Running Average Summary

Vector Check

- Where are we?
 - Looked at some Python implementations of the running average and simple moving average
 - Ran some simulations of both constant and dynamic scalar processes and examine how the two types of averages performed
- What's next?
 - In the next video we will look at the prediction part of the algorithm

Running Average Summary

Vector Check

- Where are we?
 - Looked at some Python implementations of the running average and simple moving average
 - Ran some simulations of both constant and dynamic scalar processes and examine how the two types of averages performed
- What's next?
 - In the next video we will look at the prediction part of the algorithm
 - Especially in the case where the state of the system (process) is dynamically changing

Running Average Summary

Vector Check

- Where are we?
 - Looked at some Python implementations of the running average and simple moving average
 - Ran some simulations of both constant and dynamic scalar processes and examine how the two types of averages performed
 - What's next?
 - In the next video we will look at the prediction part of the algorithm
 - Especially in the case where the state of the system (process) is dynamically changing
-